



Running Flask on Kubernetes

Posted by [Michael Herman](#) | Last updated on January 24th, 2020 |  [vue](#) [docker](#) [flask](#) [devops](#) [kubernetes](#)

In this post, we'll first take a look at Kubernetes and container orchestration in general and then we'll walk through a step-by-step tutorial that details how to deploy a Flask-based microservice (along with Postgres and Vue.js) to a Kubernetes cluster.

This is an intermediate-level tutorial. It assumes that you have basic working knowledge of Flask and Docker. Review the [Test-Driven Development with Python, Flask, and Docker](#) course for more info on these tools.

Dependencies:

- Kubernetes v1.17.1
- Minikube v1.6.2
- Docker v19.03.5
- Docker-Compose v1.24.1

Objectives

By the end of this tutorial, you will be able to:

1. Explain what container orchestration is and why you may need to use an orchestration tool
2. Discuss the pros and cons of using Kubernetes over other orchestration tools like Docker Swarm and Elastic Container Service (ECS)
3. Explain the following Kubernetes primitives: Node, Pod, Service, Label, Deployment, Ingress, and Volume
4. Spin up a Python-based microservice locally with Docker Compose
5. Configure a Kubernetes cluster to run locally with Minikube
6. Set up a volume to hold Postgres data within a Kubernetes cluster
7. Use Kubernetes Secrets to manage sensitive information
8. Run Flask, Gunicorn, Postgres, and Vue on Kubernetes
9. Expose Flask and Vue to external users via an Ingress

What is Container Orchestration?

Feedback

As you move from deploying containers on a single machine to deploying them across a number of machines, you'll need an orchestration tool to manage (and automate) the arrangement, coordination, and availability of the containers across the entire system.

Orchestration tools help with:

- 1. Cross-server container communication
- 2. Horizontal scaling
- 3. Service discovery
- 4. Load balancing
- 5. Security/TLS
- 6. Zero-downtime deploys
- 7. Rollbacks
- 8. Logging
- 9. Monitoring

This is where [Kubernetes](#) fits in along with a number of other orchestration tools -- like [Docker Swarm](#), [ECS](#), [Mesos](#), and [Nomad](#).

Which one should you use?

- use *Kubernetes* if you need to manage large, complex clusters
- use *Docker Swarm* if you are just getting started and/or need to manage small to medium-sized clusters
- use *ECS* if you're already using a number of AWS services

Tool	Pros	Cons
Kubernetes	large community, flexible, most features, hip	complex setup, high learning curve, hip
Docker Swarm	easy to set up, perfect for smaller clusters	limited by the Docker API
ECS	fully-managed service, integrated with AWS	vendor lock-in

There's also a number of [managed](#) Kubernetes services on the market:

- 1. [Google Kubernetes Engine](#) (GKE)
- 2. [Elastic Kubernetes Service](#) (EKS)
- 3. [Azure Kubernetes Service](#) (AKS)
- 4. [DigitalOcean Kubernetes](#)

For more, review the [Choosing the Right Containerization and Cluster Management Tool](#) blog post.

Kubernetes Concepts

Before diving in, let's look at some of the basic building blocks that you have to work with from the [Kubernetes API](#):

- 1. A **Node** is a worker machine provisioned to run Kubernetes. Each Node is managed by the Kubernetes master.
- 2. A **Pod** is a logical, tightly-coupled group of application containers that run on a Node. Containers in a Pod are deployed together and share resources (like data volumes and network addresses). Multiple Pods can run on a single Node.
- 3. A **Service** is a logical set of Pods that perform a similar function. It enables load balancing and service discovery. It's an abstraction layer over the Pods; Pods are meant to be ephemeral while services are much more persistent.
- 4. **Deployments** are used to describe the desired state of Kubernetes. They dictate how Pods are created, deployed, and replicated.
- 5. **Labels** are key/value pairs that are attached to resources (like Pods) which are used to organize related resources. You can think of them like CSS selectors. For example:
 - *Environment* - `dev`, `test`, `prod`
 - *App version* - `beta`, `1.2.1`
 - *Type* - `client`, `server`, `db`
- 6. **Ingress** is a set of routing rules used to control the external access to Services based on the request host or path.
- 7. **Volumes** are used to persist data beyond the life of a container. They are especially important for stateful applications like Redis and Postgres.

Feedback

- A [PersistentVolume](#) defines a storage volume independent of the normal Pod-lifecycle. It's managed outside of the particular Pod that it resides in.
- A [PersistentVolumeClaim](#) is a request to use the PersistentVolume by a user.

For more, review the [Learn Kubernetes Basics](#) tutorial as well as the [Kubernetes Concepts](#) slides from the [Scaling Flask with Kubernetes](#) talk.

Project Setup

Clone down the [flask-vue-kubernetes](#) repo, and then build the images and spin up the containers:

```
$ git clone https://github.com/testdrivenio/flask-vue-kubernetes
$ cd flask-vue-kubernetes
$ docker-compose up -d --build
```

Create and seed the database `books` table:

```
$ docker-compose exec server python manage.py recreate_db
$ docker-compose exec server python manage.py seed_db
```

Test out the following server-side endpoints in your browser of choice.

<http://localhost:5001/books/ping>

```
{
  "container_id": "dee114fa81ea",
  "message": "pong!",
  "status": "success"
}
```

`container_id` is the id of the Docker container the app is running in.

```
$ docker ps --filter name=flask-vue-kubernetes_server --format "{{.ID}}"

dee114fa81ea
```

<http://localhost:5001/books>:

```
{
  "books": [{
    "author": "J. K. Rowling",
    "id": 2,
    "read": false,
    "title": "Harry Potter and the Philosopher's Stone"
  }, {
    "author": "Dr. Seuss",
    "id": 3,
    "read": true,
    "title": "Green Eggs and Ham"
  }, {
    "author": "Jack Kerouac",
    "id": 1,
    "read": false,
    "title": "On the Road"
  }],
  "container_id": "dee114fa81ea",
  "status": "success"
}
```

Navigate to <http://localhost:8080>. Make sure the basic CRUD functionality works as expected:



Books

Add Book



Title	Author	Read?		
On the Road	Jack Kerouac	Yes	Update	Delete
Harry Potter and the Philosopher's Stone	J. K. Rowling	No	Update	Delete
Green Eggs and Ham	Dr. Seuss	Yes	Update	Delete
1Q84	Haruki Murakami	Yes	Update	Delete

Take a quick look at the code before moving on:

Feedback

```
├─ .gitignore
├─ README.md
├─ deploy.sh
├─ docker-compose.yml
├─ kubernetes
│   ├─ flask-deployment.yml
│   ├─ flask-service.yml
│   ├─ minikube-ingress.yml
│   ├─ persistent-volume-claim.yml
│   ├─ persistent-volume.yml
│   ├─ postgres-deployment.yml
│   ├─ postgres-service.yml
│   ├─ secret.yml
│   ├─ vue-deployment.yml
│   └─ vue-service.yml
├─ services
│   └─ client
│       ├─ .babelrc
│       ├─ .editorconfig
│       ├─ .eslintignore
│       ├─ .eslintrc.js
│       ├─ .postcssrc.js
│       ├─ Dockerfile
│       ├─ Dockerfile-minikube
│       ├─ README.md
│       ├─ build
│       ├─ config
│       │   ├─ dev.env.js
│       │   ├─ index.js
│       │   └─ prod.env.js
│       ├─ index.html
│       ├─ package-lock.json
│       ├─ package.json
│       ├─ src
│       │   ├─ App.vue
│       │   ├─ assets
│       │   │   └─ logo.png
│       │   ├─ components
│       │   │   ├─ Alert.vue
│       │   │   ├─ Books.vue
│       │   │   └─ HelloWorld.vue
│       │   └─ Ping.vue
│       ├─ main.js
│       └─ router
│           └─ index.js
│       └─ static
│           └─ .gitkeep
├─ db
│   ├─ create.sql
│   └─ dockerfile
├─ server
│   ├─ .dockerignore
│   ├─ dockerfile
│   ├─ entrypoint.sh
│   ├─ manage.py
│   ├─ project
│   │   ├─ __init__.py
│   │   ├─ api
│   │   │   └─ __init__.py
│   │   ├─ books.py
│   │   └─ models.py
│   └─ config.py
└─ requirements.txt
```

Want to learn how to build this project? Check out the [Developing a Single Page App with Flask and Vue.js](#) blog post.

Minikube

[Minikube](#) is a tool which allows developers to use and run a Kubernetes cluster locally. It's a great way to quickly get a cluster up and running so you can start interacting with the Kubernetes API.

Follow the official [quickstart](#) guide to get Minikube installed along with:

- 1. A [Hypervisor](#) (like [VirtualBox](#) or [HyperKit](#)) to manage virtual machines

Feedback

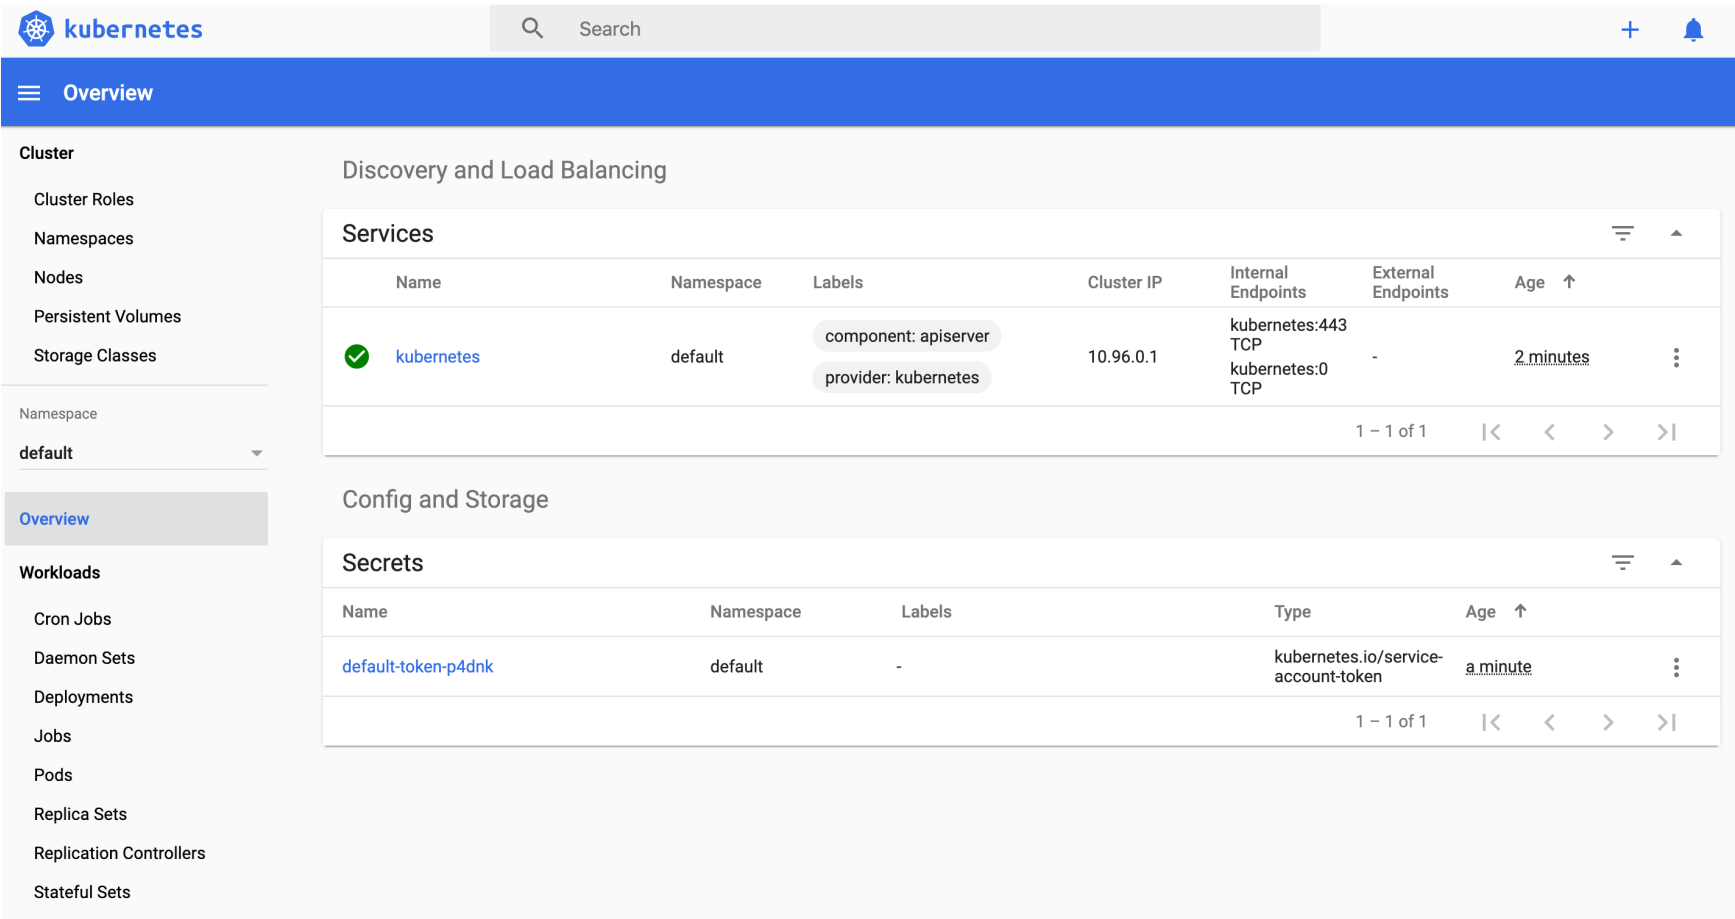
2. [Kubectl](#) to deploy and manage apps on Kubernetes

If you're on a Mac, we recommend installing Kubectl and Minikube with [Homebrew](#):

```
$ brew update
$ brew install kubectl
$ brew install minikube
```

Then, start the cluster and pull up the Minikube [dashboard](#):

```
$ minikube config set vm-driver hyperkit
$ minikube start
$ minikube dashboard
```



It's worth noting that the config files will be located in the `~/.kube` directory while all the virtual machine bits will be in the `~/.minikube` directory.

Now we can start creating objects via the Kubernetes API.

If you run into problems with Minikube, it's often best to remove it completely and start over.

For example:

```
$ minikube stop; minikube delete
$ rm /usr/local/bin/minikube
$ rm -rf ~/.minikube
# re-download minikube
$ minikube start
```

Creating Objects

To create a new [object](#) in Kubernetes, you must provide a "spec" that describes its desired state.

Example:

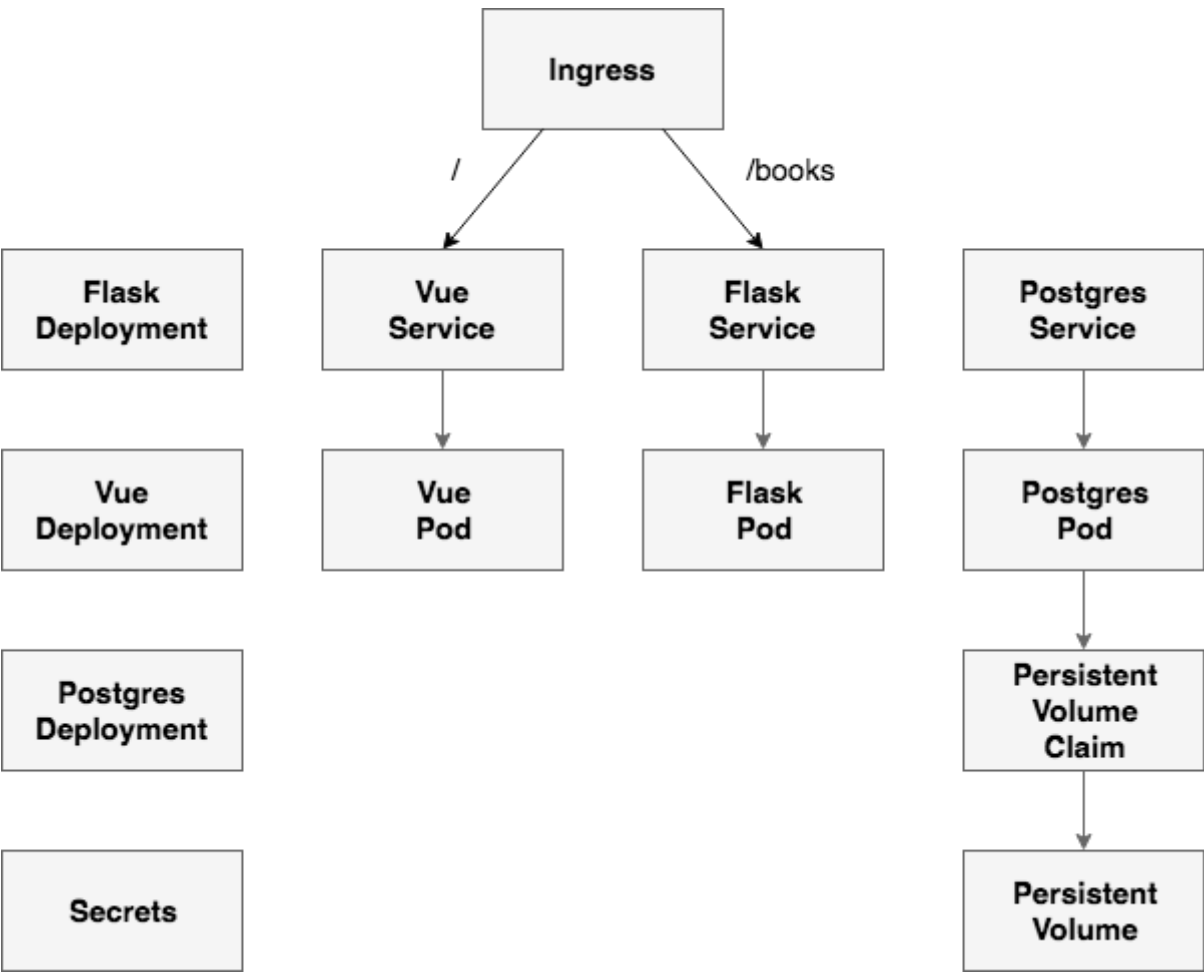
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: flask
    spec:
      containers:
        - name: flask
          image: mjhea0/flask-kubernetes:latest
          ports:
            - containerPort: 5000
```

Required Fields:

- 1. `apiVersion` - [Kubernetes API](#) version
- 2. `kind` - the type of object you want to create
- 3. `metadata` - info about the object so that it can be uniquely identified
- 4. `spec` - desired state of the object

In the above example, this spec will create a new Deployment for a Flask app with a single replica (Pod). Take note of the `containers` section. Here, we specified the Docker image along with the container port the application will run on.

In order to run our app, we'll need to set up the following objects:



Volume

Again, since containers are ephemeral, we need to configure a volume, via a [PersistentVolume](#) and a [PersistentVolumeClaim](#), to store the Postgres data outside of the Pod.

Take note of the YAML file in `kubernetes/persistent-volume.yml`:


```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: postgres-pv
  labels:
    type: local
spec:
  capacity:
    storage: 2Gi
  storageClassName: standard
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/data/postgres-pv"
```

This configuration will create a [hostPath](#) PersistentVolume at "/data/postgres-pv" within the Node. The size of the volume is 2 gibibytes with an access mode of [ReadWriteOnce](#), which means that the volume can be mounted as read-write by a single node.

It's worth noting that Kubernetes only supports using a hostPath on a single-node cluster.

Create the volume:

```
$ kubectl apply -f ./kubernetes/persistent-volume.yml
```

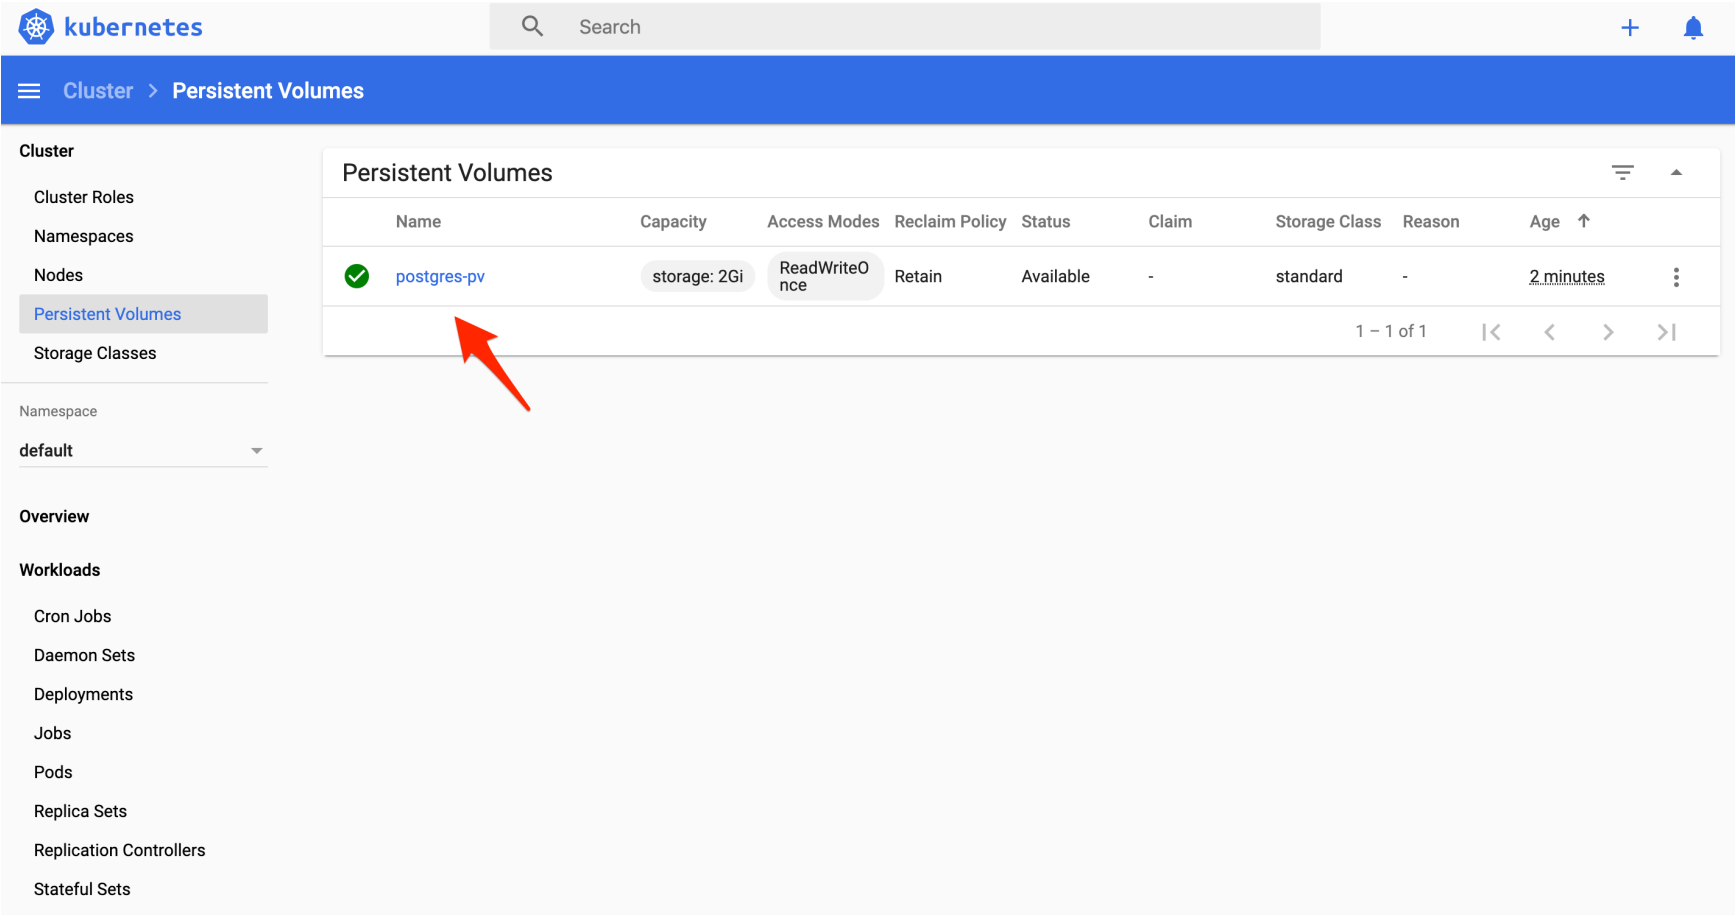
View details:

```
$ kubectl get pv
```

You should see:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
postgres-pv	2Gi	RWO	Retain	Available		standard		14s

You should also see this object in the dashboard:



kubernetes/persistent-volume-claim.yml:


```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: postgres-pvc
  labels:
    type: local
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
  volumeName: postgres-pv
  storageClassName: standard
```

Create the volume claim:

```
$ kubectl apply -f ./kubernetes/persistent-volume-claim.yml
```

View details:

```
$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
postgres-pvc	Bound	postgres-pv	2Gi	RWO	standard	15s

kubernetes

Search

Config and Storage > Persistent Volume Claims > postgres-pvc

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Stateful Sets

Discovery and Load Balancing

Ingresses

Services

Config and Storage

Config Maps

Persistent Volume Claims

Secrets

Custom Resource Definitions

Settings

About

Metadata

Name	Namespace	Creation time	Age	UID
postgres-pvc	default	Jan 23, 2020	a minute	e08ed6b0-7e2b-4251-92db-516b59948d53

Labels

type: local

Annotations

[kubectl.kubernetes.io/last-applied-configuration](#)pv.kubernetes.io/bind-completed: yes

Resource information

Status	Storage Class
Bound	standard

Capacity

storage: 2Gi

Access Modes

ReadWriteOnce

Secrets

[Secrets](#) are used to handle sensitive info such as passwords, API tokens, and SSH keys. We'll set up a Secret to store our Postgres database credentials.

kubernetes/secret.yml:

```
apiVersion: v1
kind: Secret
metadata:
  name: postgres-credentials
type: Opaque
data:
  user: c2FtcGx1
  password: cGx1YXNlY2hhbmdlbWU=
```

The `user` and `password` fields are base64 encoded strings ([security via obscurity](#)):

Feedback

<https://testdriven.io/blog/running-flask-on-kubernetes/>

9/21

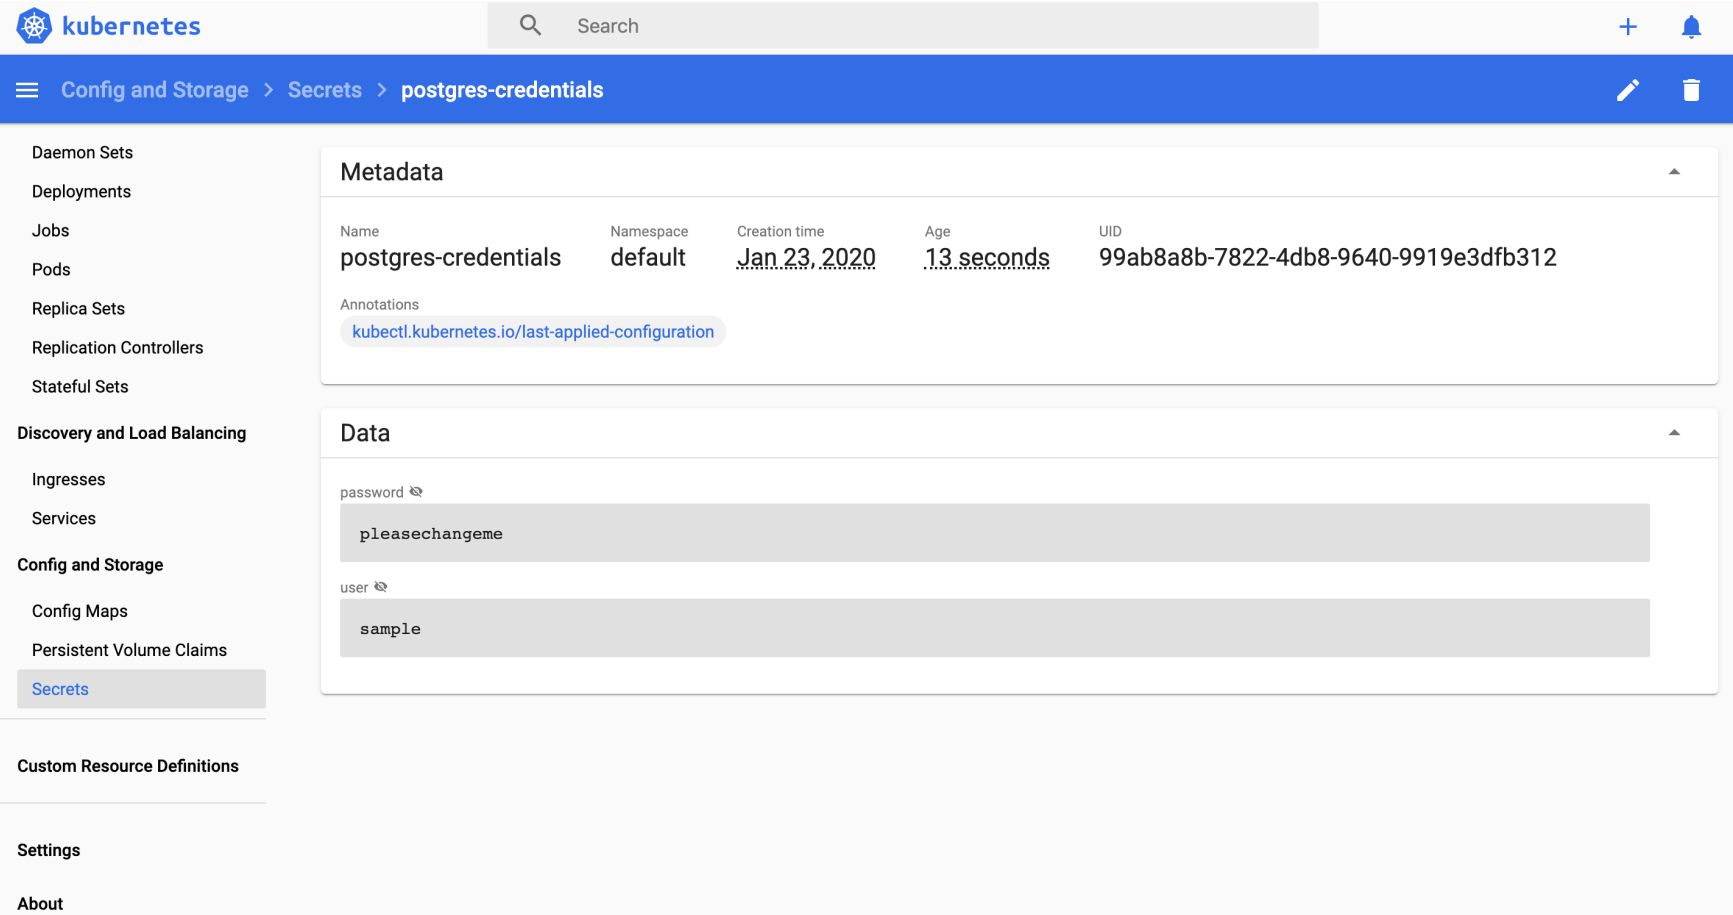
```
$ echo -n "pleasechangeme" | base64
cGx1YXNlY2hhbmdlbWU=

$ echo -n "sample" | base64
c2FtcGx1
```

Keep in mind that any user with access to the cluster will be able to read the values in plaintext. Take a look at [Vault](#) if you want to encrypt secrets in transit and at rest.

Add the Secrets object:

```
$ kubectl apply -f ./kubernetes/secret.yml
```



Postgres

With the volume and database credentials set up in the cluster, we can now configure the Postgres database itself.

kubernetes/postgres-deployment.yml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres
  labels:
    name: database
spec:
  replicas: 1
  selector:
    matchLabels:
      service: postgres
  template:
    metadata:
      labels:
        service: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:12-alpine
          env:
            - name: POSTGRES_USER
              valueFrom:
                secretKeyRef:
                  name: postgres-credentials
                  key: user
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-credentials
                  key: password
          volumeMounts:
            - name: postgres-volume-mount
              mountPath: /var/lib/postgresql/data
      volumes:
        - name: postgres-volume-mount
          persistentVolumeClaim:
            claimName: postgres-pvc
      restartPolicy: Always
```

What's happening here?

- 1. `metadata`
 - The `name` field defines the Deployment name - `postgres`
 - `labels` define the labels for the Deployment - `name: database`
- 2. `spec`
 - `replicas` define the number of Pods to run - `1`
 - `selector` defines how the Deployment finds which Pods to manage
 - `template`
 - `metadata`
 - `labels` indicate which labels should be assigned to the Pod - `service: postgres`
 - `spec`
 - `containers` define the containers associated with each Pod
 - `volumes` define the volume claim - `postgres-volume-mount`
 - `restartPolicy` defines the [restart policy](#) - `Always`

Further, the Pod name is `postgres` and the image is `postgres:12.1-alpine`, which will be pulled from Docker Hub. The database credentials, from the Secret object, are passed in as well.

Finally, when applied, the volume claim will be mounted into the Pod. The claim is mounted to `"/var/lib/postgresql/data"` - the default location - while the data will be stored in the PersistentVolume, `"/data/postgres-pv"`.

Create the Deployment:

```
$ kubectl create -f ./kubernetes/postgres-deployment.yml
```

Status:

```
$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
postgres	1/1	1	1	12s

Feedback

kubernetes/postgres-service.yml:

```
apiVersion: v1
kind: Service
metadata:
  name: postgres
  labels:
    service: postgres
spec:
  selector:
    service: postgres
  type: ClusterIP
  ports:
    - port: 5432
```

What's happening here?

1. metadata
- The name field defines the Service name - postgres

◦ labels define the labels for the Service - name: database
2. spec
- selector defines the Pod label and value that the Service applies to - service: postgres

◦ type defines the type of Service - ClusterIP

◦ ports

▪ port defines the port exposed to the cluster

Take a moment to go back to the Deployment spec. How does the selector in the Service relate back to the Deployment?

Since the Service type is ClusterIP, it's not exposed externally, so it's *only* accessible from within the cluster by other objects.

Create the service:

```
$ kubectl create -f ./kubernetes/postgres-service.yml
```

Create the books database, using the Pod name:

```
$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
postgres-95566f9-xs2cf             1/1     Running   0           93s

$ kubectl exec postgres-95566f9-xs2cf --stdin --tty -- createdb -U sample books
```

Verify the creation:

```
$ kubectl exec postgres-95566f9-xs2cf --stdin --tty -- psql -U sample

psql (12.1)
Type "help" for help.

sample=# \l

               List of databases
  Name  | Owner   | Encoding | Collate  |  Ctype  | Access privileges
-----+-----+-----+-----+-----+-----
 books  | sample  | UTF8     | en_US.utf8 | en_US.utf8 | 
 postgres | postgres | UTF8     | en_US.utf8 | en_US.utf8 | 
 sample | postgres | UTF8     | en_US.utf8 | en_US.utf8 | 
 template0 | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =c/postgres +
         |         |         |         |         | postgres=CTc/postgres
 template1 | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =c/postgres +
         |         |         |         |         | postgres=CTc/postgres
(5 rows)

sample=#
```

You can also get the Pod name via:

Feedback

```
$ kubectl get pod -l service=postgres -o jsonpath="{.items[0].metadata.name}"
```

Assign the value to a variable and then create the database:

```
$ POD_NAME=$(kubectl get pod -l service=postgres -o jsonpath="{.items[0].metadata.name}")
$ kubectl exec $POD_NAME --stdin --tty -- createdb -U sample books
```

Flask

Take a moment to review the Flask project structure along with the *dockerfile* and the *entrypoint.sh* files:

- 1. "services/server"
- 2. *services/server/dockerfile*
- 3. *services/server/entrypoint.sh*

kubernetes/flask-deployment.yml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask
  labels:
    name: flask
spec:
  replicas: 1
  selector:
    matchLabels:
      app: flask
  template:
    metadata:
      labels:
        app: flask
    spec:
      containers:
        - name: flask
          image: mjhea0/flask-kubernetes:latest
          env:
            - name: FLASK_ENV
              value: "development"
            - name: APP_SETTINGS
              value: "project.config.DevelopmentConfig"
            - name: POSTGRES_USER
              valueFrom:
                secretKeyRef:
                  name: postgres-credentials
                  key: user
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-credentials
                  key: password
          restartPolicy: Always
```

This should look similar to the Postgres Deployment spec. The big difference is that you can either use my pre-built and pre-pushed image on [Docker Hub](#), `mjhea0/flask-kubernetes`, or build and push your own.

For example:

```
$ docker build -t <YOUR_DOCKER_HUB_NAME>/flask-kubernetes ./services/server
$ docker push <YOUR_DOCKER_HUB_NAME>/flask-kubernetes
```

If you use your own, make sure you replace `mjhea0` with your Docker Hub name in *kubernetes/flask-deployment.yml* as well.

Alternatively, if don't want to push the image to a Docker registry, after you build the image locally, you can set the `image-pull-policy` flag to `Never` to always use the local image.

Create the Deployment:

```
$ kubectl create -f ./kubernetes/flask-deployment.yml
```

kubernetes

Search

+

Workloads > Deployments > flask

Cluster

Cluster Roles

Namespaces

Nodes

Persistent Volumes

Storage Classes

Namespace

default

Overview

Workloads

Cron Jobs

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Stateful Sets

Metadata

Name	Namespace	Creation time	Age	UID
flask	default	Jan 23, 2020	a minute	9cb136f0-dc45-466f-b9cd-c7f7903e0e65

Labels

name: flask

Annotations

deployment.kubernetes.io/revision: 1

Resource information

Strategy	Min ready seconds	Revision history limit
RollingUpdate	0	10

Selector

app: flask

Rolling update strategy

Max surge	Max unavailable
25%	25%

This will immediately spin up a new Pod:

kubernetes

Search

+

Workloads > Pods > flask-66988cb97d-n88b4

Cluster

Cluster Roles

Namespaces

Nodes

Persistent Volumes

Storage Classes

Namespace

default

Overview

Workloads

Cron Jobs

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Stateful Sets

Metadata

Name	Namespace	Creation time	Age	UID
flask-66988cb97d-n88b4	default	Jan 23, 2020	2 minutes	a75be219-cbe4-47ad-9d57-e47d2f93f713

Labels

app: flask

pod-template-hash: 66988cb97d

Resource information

Node	Status	IP	QoS Class	Restarts
minikube	Running	172.17.0.7	BestEffort	0

Conditions

Type	Status	Last probe time	Last transition time	Reason	Message
Initialized	True	::	2 minutes	-	-
Ready	True	::	a minute	-	-
ContainersReady	True	::	a minute	-	-
PodScheduled	True	::	2 minutes	-	-

kubernetes/flask-service.yml:

```
apiVersion: v1
kind: Service
metadata:
  name: flask
  labels:
    service: flask
spec:
  selector:
    app: flask
  ports:
    - port: 5000
      targetPort: 5000
```

Curious about the `targetPort` and how it relates to the `port`? Review the official [Services](#) guide.

Feedback

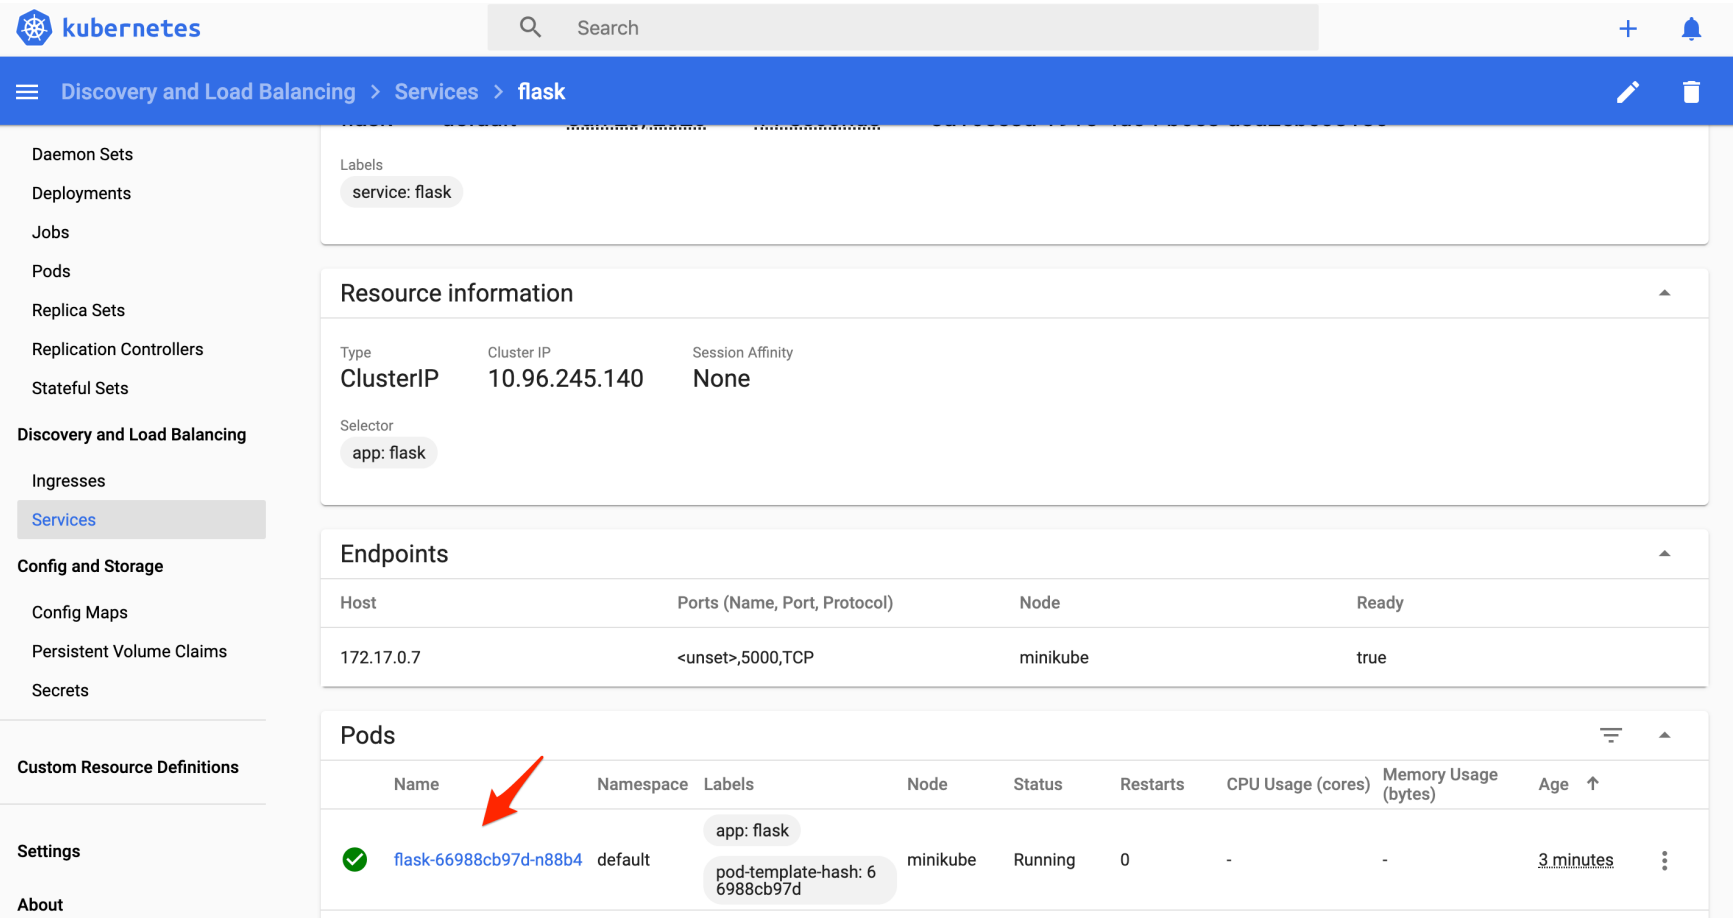
https://testdriven.io/blog/running-flask-on-kubernetes/

14/21

Create the service:

```
$ kubectl create -f ./kubernetes/flask-service.yml
```

Make sure the Pod is associated with the Service:



Apply the migrations and seed the database:

```
$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
flask-66988cb97d-n88b4             1/1     Running   0           21m
postgres-95566f9-xs2cf             1/1     Running   0           36m
```

```
$ kubectl exec flask-66988cb97d-n88b4 --stdin --tty -- python manage.py recreate_db
$ kubectl exec flask-66988cb97d-n88b4 --stdin --tty -- python manage.py seed_db
```

Verify:

```
$ kubectl exec postgres-95566f9-xs2cf --stdin --tty -- psql -U sample

psql (12.1)
Type "help" for help.

sample=# \c books
You are now connected to database "books" as user "sample".
books=# select * from books;

 id | title                                     | author      | read
----+-----+-----+-----
  1 | On the Road                             | Jack Kerouac | t
  2 | Harry Potter and the Philosopher's Stone | J. K. Rowling | f
  3 | Green Eggs and Ham                      | Dr. Seuss    | t
(3 rows)
```

Ingress

To enable traffic to access the Flask API inside the cluster, you can use either a NodePort, LoadBalancer, or Ingress:

- 1. A [NodePort](#) exposes a Service on an open port on the Node.
- 2. As the name implies, a [LoadBalancer](#) creates an external load balancer that points to a Service in the cluster.
- 3. Unlike the previous two methods, an [Ingress](#) is not a type of Service; instead, it sits on top of the Services as an entry point into the cluster.

For more, review the official [Publishing Services](#) guide.

kubernetes/minikube-ingress.yml:

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: minikube-ingress
  annotations:
spec:
  rules:
  - host: hello.world
    http:
      paths:
      - path: /
        backend:
          serviceName: vue
          servicePort: 8080
      - path: /books
        backend:
          serviceName: flask
          servicePort: 5000
```

Here, we defined the following HTTP rules:

1. `/` - routes requests to the Vue Service (which we still need to set up)
2. `/books` - routes requests to the Flask Service

Enable the Ingress [addon](#):

```
$ minikube addons enable ingress
```

Create the Ingress object:

```
$ kubectl apply -f ./kubernetes/minikube-ingress.yml
```

Next, you need to update your `/etc/hosts` file to route requests from the host we defined, `hello.world`, to the Minikube instance.

Add an entry to `/etc/hosts`:

```
$ echo "$(minikube ip) hello.world" | sudo tee -a /etc/hosts
```

Try it out:

```
http://hello.world/books/ping
```

```
{
  "container_id": "flask-66988cb97d-n88b4",
  "message": "pong!", "status":
  "success"
}
```

```
http://hello.world/books
```

```
{
  "books": [{
    "author": "Jack Kerouac",
    "id": 1,
    "read": true,
    "title": "On the Road"
  }, {
    "author": "J. K. Rowling",
    "id": 2,
    "read": false,
    "title": "Harry Potter and the Philosopher's Stone"
  }, {
    "author": "Dr. Seuss",
    "id": 3,
    "read": true,
    "title": "Green Eggs and Ham"
  }],
  "container_id": "flask-66988cb97d-n88b4",
  "status": "success"
}
```

Vue

Moving right along, review the Vue project along with the associated Dockerfiles:

- 1. "services/client"
- 2. ./services/client/Dockerfile
- 3. ./services/client/Dockerfile-minikube

kubernetes/vue-deployment.yml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: vue
  labels:
    name: vue
spec:
  replicas: 1
  selector:
    matchLabels:
      app: vue
  template:
    metadata:
      labels:
        app: vue
    spec:
      containers:
        - name: vue
          image: mjhea0/vue-kubernetes:latest
          restartPolicy: Always
```

Again, either use my image or build and push your own image to Docker Hub:

```
$ docker build -t <YOUR_DOCKERHUB_NAME>/vue-kubernetes ./services/client \
  -f ./services/client/Dockerfile-minikube
$ docker push <YOUR_DOCKERHUB_NAME>/vue-kubernetes
```

Create the Deployment:

```
$ kubectl create -f ./kubernetes/vue-deployment.yml
```

Verify that a Pod was created along with the Deployment:

```
$ kubectl get deployments vue

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
vue       1/1     1            1           40s

$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
flask-66988cb97d-n88b4             1/1     Running   0          37m
postgres-95566f9-xs2cf             1/1     Running   0          71m
vue-cd9d7d445-xl7wd                1/1     Running   0          2m32s
```

How would you verify that the Pod and Deployment were created successfully in the dashboard?

kubernetes/vue-service.yml:

```
apiVersion: v1
kind: Service
metadata:
  name: vue
  labels:
    service: vue
  name: vue
spec:
  selector:
    app: vue
  ports:
    - port: 8080
      targetPort: 8080
```

Create the service:

```
$ kubectl create -f ./kubernetes/vue-service.yml
```

Ensure <http://hello.world/> works as expected.

Books

Book updated!

Add Book

Title	Author	Read?		
Harry Potter and the Philosopher's Stone	J. K. Rowling	No	Update	Delete
Green Eggs and Ham	Dr. Seuss	Yes	Update	Delete
On the Road	Jack Kerouac	No	Update	Delete

Scaling

Kubernetes makes it easy to scale, adding additional Pods as necessary, when the traffic load becomes too much for a single Pod to handle.

For example, let's add another Flask Pod to the cluster:

Feedback

```
$ kubectl scale deployment flask --replicas=2
```

Confirm:

```
$ kubectl get deployments flask
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
flask	2/2	2	2	11m

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
flask-66988cb97d-hqpbh	1/1	Running	0	27s	172.17.0.10	minikube	<none>		<none>	
flask-66988cb97d-n88b4	1/1	Running	0	39m	172.17.0.7	minikube	<none>		<none>	
postgres-95566f9-xs2cf	1/1	Running	0	74m	172.17.0.6	minikube	<none>		<none>	
vue-cd9d7d445-xl7wd	1/1	Running	0	5m18s	172.17.0.9	minikube	<none>		<none>	

Make a few requests to the service:

```
$ for ((i=1;i<=10;i++)); do curl http://hello.world/books/ping; done
```

You should see different `container_id`s being returned, indicating that requests are being routed appropriately via a round robin algorithm between the two replicas:

```
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-n88b4","message":"pong!","status":"success"}
{"container_id":"flask-66988cb97d-hqpbh","message":"pong!","status":"success"}
```

What happens if you scale down as traffic is hitting the cluster? Open two terminal windows and test this on your on. You should see traffic being re-routed appropriately. Try it again, but this time scale up.

Helpful Commands

Command	Explanation
<code>minikube start</code>	Starts a local Kubernetes cluster
<code>minikube ip</code>	Displays the IP address of the cluster
<code>minikube dashboard</code>	Opens the Kubernetes dashboard in your browser
<code>kubectl version</code>	Displays the Kubectl version
<code>kubectl cluster-info</code>	Displays the cluster info
<code>kubectl get nodes</code>	Lists the Nodes
<code>kubectl get pods</code>	Lists the Pods
<code>kubectl get deployments</code>	Lists the Deployments
<code>kubectl get services</code>	Lists the Services
<code>minikube stop</code>	Stops a local Kubernetes cluster
<code>minikube delete</code>	Removes a local Kubernetes cluster

Check out the [Kubernetes Cheatsheet](#) for more commands.

Automation Script

Ready to put everything together?

Take a look at the `deploy.sh` script in the project root. This script:

- 1. Creates a PersistentVolume and a PersistentVolumeClaim
- 2. Adds the database credentials via Kubernetes Secrets
- 3. Creates the Postgres Deployment and Service
- 4. Creates the Flask Deployment and Service
- 5. Enables Ingress
- 6. Applies the Ingress rules
- 7. Creates the Vue Deployment and Service

```
#!/bin/bash

echo "Creating the volume..."

kubectl apply -f ./kubernetes/persistent-volume.yml
kubectl apply -f ./kubernetes/persistent-volume-claim.yml

echo "Creating the database credentials..."

kubectl apply -f ./kubernetes/secret.yml

echo "Creating the postgres deployment and service..."

kubectl create -f ./kubernetes/postgres-deployment.yml
kubectl create -f ./kubernetes/postgres-service.yml
POD_NAME=$(kubectl get pod -l service=postgres -o jsonpath="{.items[0].metadata.name}")
kubectl exec $POD_NAME --stdin --tty -- createdb -U sample books

echo "Creating the flask deployment and service..."

kubectl create -f ./kubernetes/flask-deployment.yml
kubectl create -f ./kubernetes/flask-service.yml
FLASK_POD_NAME=$(kubectl get pod -l app=flask -o jsonpath="{.items[0].metadata.name}")
kubectl exec $FLASK_POD_NAME --stdin --tty -- python manage.py recreate_db
kubectl exec $FLASK_POD_NAME --stdin --tty -- python manage.py seed_db

echo "Adding the ingress..."

minikube addons enable ingress
kubectl apply -f ./kubernetes/minikube-ingress.yml

echo "Creating the vue deployment and service..."

kubectl create -f ./kubernetes/vue-deployment.yml
kubectl create -f ./kubernetes/vue-service.yml
```

Try it out!

```
$ sh deploy.sh
```

Once done, create the `books` database, apply the migrations, and seed the database:

```
$ POD_NAME=$(kubectl get pod -l service=postgres -o jsonpath="{.items[0].metadata.name}")
$ kubectl exec $POD_NAME --stdin --tty -- createdb -U sample books

$ FLASK_POD_NAME=$(kubectl get pod -l app=flask -o jsonpath="{.items[0].metadata.name}")
$ kubectl exec $FLASK_POD_NAME --stdin --tty -- python manage.py recreate_db
$ kubectl exec $FLASK_POD_NAME --stdin --tty -- python manage.py seed_db
```

Feedback

Update `/etc/hosts`, and then test it out in the browser.

Conclusion

In this post we looked at how to run a Flask-based microservice on Kubernetes.

At this point, you should have a basic understanding of how Kubernetes works and be able to deploy a cluster with an app running on it.

Additional Resources:

- 1. [Learn Kubernetes Basics](#)
- 2. [Configuration Best Practices](#)
- 3. [Scaling Flask with Kubernetes](#)
- 4. [Running Flask on Docker Swarm](#) (compare and contrast running Flask on Docker Swarm vs. Kubernetes)
- 5. [Deploying a Node App to Google Cloud with Kubernetes](#)

You can find the code in the [flask-vue-kubernetes](#) repo on GitHub.

 [vue](#) [docker](#) [flask](#) [devops](#) [kubernetes](#)

Share this tutorial

 Twitter

 Reddit

 Hacker News

 Facebook

Join our mailing list to be notified about course updates and new tutorials.

Enter your email...

Subscribe

Learn Vue by Building and Deploying a CRUD App

Get the full course. Learn Vue by building and testing a full-featured web application.

View the Course